

Optional Lab: Linear Regression using Scikit-Learn

There is an open-source, commercially usable machine learning toolkit called [scikit-learn](#). This toolkit contains implementations of many of the algorithms that you will work with in this course.

Goals

In this lab you will:

- Utilize scikit-learn to implement linear regression using a close form solution based on the normal equation

Tools

You will utilize functions from scikit-learn as well as matplotlib and NumPy.

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from lab_utils_multi import load_house_data
plt.style.use('./deeplearning.mplstyle')
np.set_printoptions(precision=2)
```

Linear Regression, closed-form solution

Scikit-learn has the [linear regression model](#) which implements a closed-form linear regression.

Let's use the data from the early labs - a house with 1000 square feet sold for \$300,000 and a house with 2000 square feet sold for \$500,000.

Size (1000 sqft)	Price (1000s of dollars)
1	300
2	500

Load the data set

```
In [2]: X_train = np.array([1.0, 2.0]) #features
y_train = np.array([300, 500]) #target value
```

Create and fit the model

The code below performs regression using scikit-learn. The first step creates a regression object. The second step utilizes one of the methods associated with the object, `fit`. This performs regression, fitting the parameters to the input data. The toolkit expects a two-dimensional X matrix.

```
In [3]: linear_model = LinearRegression()  
#X must be a 2-D Matrix  
linear_model.fit(X_train.reshape(-1, 1), y_train)
```

```
Out[3]: ▼ LinearRegression  
LinearRegression()
```

View Parameters

The **w** and **b** parameters are referred to as 'coefficients' and 'intercept' in scikit-learn.

```
In [4]: b = linear_model.intercept_  
w = linear_model.coef_  
print(f"w = {w:}, b = {b:0.2f}")  
print(f"'manual' prediction: f_wb = wx+b : {1200*w + b}")  
  
w = [200.], b = 100.00  
'manual' prediction: f_wb = wx+b : [240100.]
```

Make Predictions

Calling the `predict` function generates predictions.

```
In [5]: y_pred = linear_model.predict(X_train.reshape(-1, 1))  
  
print("Prediction on training set:", y_pred)  
  
X_test = np.array([[1200]])  
print(f"Prediction for 1200 sqft house: ${linear_model.predict(X_test)[0]:0.2f}")  
  
Prediction on training set: [300. 500.]  
Prediction for 1200 sqft house: $240100.00
```

Second Example

The second example is from an earlier lab with multiple features. The final parameter values and predictions are very close to the results from the un-normalized 'long-run' from that lab. That un-normalized run took hours to produce results, while this is nearly instantaneous. The closed-form solution work well on smaller data sets such as these but can be computationally demanding on larger data sets.

The closed-form solution does not require normalization.

```
In [6]: # load the dataset  
X_train, y_train = load_house_data()  
X_features = ['size(sqft)', 'bedrooms', 'floors', 'age']
```

```
In [7]: linear_model = LinearRegression()  
linear_model.fit(X_train, y_train)
```

```
Out[7]: ▼ LinearRegression  
LinearRegression()
```

```
In [8]: b = linear_model.intercept_  
w = linear_model.coef_  
print(f"w = {w:}, b = {b:0.2f}")
```

```
w = [ 0.27 -32.62 -67.25 -1.47], b = 220.42
```

```
In [9]: print(f"Prediction on training set:\n {linear_model.predict(X_train)[:4]}")
print(f"prediction using w,b:\n {(X_train @ w + b)[:4]}")
print(f"Target values \n {y_train[:4]}")

x_house = np.array([1200, 3, 1, 40]).reshape(-1, 4)
x_house_predict = linear_model.predict(x_house)[0]
print(f" predicted price of a house with 1200 sqft, 3 bedrooms, 1 floor, 40 years old =

Prediction on training set:
[295.18 485.98 389.52 492.15]
prediction using w,b:
[295.18 485.98 389.52 492.15]
Target values
[300.  509.8 394.  540. ]
predicted price of a house with 1200 sqft, 3 bedrooms, 1 floor, 40 years old = $318709.
09
```

Congratulations!

In this lab you:

- utilized an open-source machine learning toolkit, scikit-learn
- implemented linear regression using a close-form solution from that toolkit

```
In [ ]:
```